

# Enterprise RAG System

## Documentation

---

Complete Technical Guide

---

**Version:** 1.0.0

**Last Updated:** December 2024

**Author:** Technical Documentation Team

---

## Table of Contents

---

- 1. Executive Summary
  - 2. System Architecture
  - 3. Core Components
  - 4. Ingestion Pipeline
  - 5. Query Pipeline
  - 6. Hallucination Prevention
  - 7. API Reference
  - 8. Installation & Setup
  - 9. Configuration Reference
  - 10. Appendix
- 

## 1. Executive Summary

---

### What is RAG?

---

**Retrieval-Augmented Generation (RAG)** is an AI architecture that enhances Large Language Models (LLMs) by grounding their responses in actual documents. Instead of relying solely on

the model's training data, RAG systems:

1. **Retrieve** relevant document chunks based on the user's query
2. **Augment** the LLM prompt with this retrieved context
3. **Generate** answers strictly from the provided context

## Why This System?


---

This Enterprise RAG System is a **backend-only** implementation designed for production use. Key features:

Feature	Description
<b>No Hallucination</b>	Answers only from ingested documents
<b>Mandatory Citations</b>	Every answer includes source references
<b>Persistent Storage</b>	FAISS index and metadata survive restarts
<b>Embedding Cache</b>	SHA-256 based caching for efficiency
<b>Multi-Format Support</b>	PDF, DOCX, and Markdown documents

## Technology Stack

---




## 2. System Architecture

---

### High-Level Architecture Diagram

---



## Component Overview

---

Component	Responsibility	Key Files
FastAPI Server	HTTP endpoint handling, request validation	<code>main.py</code> , <code>api/routes/</code>
Ingestion Service	Document loading, chunking, processing	<code>ingestion/pipeline.py</code>
Embedding Service	Vector generation with OpenAI	<code>core/embeddings.py</code>
Vector Store	FAISS index management, similarity search	<code>core/vector_store.py</code>
LLM Service	Context-constrained answer generation	<code>core/llm.py</code>
Retriever Service	RAG pipeline orchestration	<code>core/retriever.py</code>

## Data Flow Overview

---



---

## 3. Core Components

---

### 3.1 Embedding Service ([core/embeddings.py](#))


---

The Embedding Service converts text into numerical vectors (embeddings) that capture semantic meaning. These vectors enable similarity search.

## Key Features

Feature	Implementation
<b>Model</b>	OpenAI <code>text-embedding-3-small</code> (1536 dimensions)
<b>Caching</b>	SHA-256 hash-based persistent cache
<b>Batching</b>	Process up to 100 texts per API call
<b>Retry Logic</b>	Exponential backoff for rate limits

## Architecture Diagram



## Cache Benefits

- Idempotent Indexing:** Same text always produces same embedding
- Cost Reduction:** Avoid redundant API calls
- Faster Rebuilds:** Cached embeddings persist across restarts

## Key Methods

Method	Purpose
<code>get_embedding(text)</code>	Single text embedding (cached)
<code>get_embeddings_batch(texts)</code>	Batch embedding with cache
<code>get_query_embedding(query)</code>	Query embedding (uncached)
<code>clear_cache()</code>	Clear all cached embeddings

## 3.2 LLM Service ( `core/llm.py` )

The LLM Service generates answers strictly from provided context. It is the core component ensuring **no hallucination**.

### Anti-Hallucination System Prompt

The LLM receives this strict system prompt:

```
ABSOLUTE RULES - VIOLATION IS FORBIDDEN:  
1. You may ONLY use information explicitly stated in the provided CONTEXT.  
2. You must NEVER use your training data, prior knowledge, or make assumptions.  
3. If the CONTEXT does not contain enough information to answer, you MUST respond:  
    "I don't know based on the provided documents."  
4. If the CONTEXT is empty or completely irrelevant, you MUST respond:  
    "Answer not found in documents."  
5. Always cite which source(s) you used with [Source N] notation.
```

## Configuration

Parameter	Default	Description
<code>LLM_MODEL</code>	<code>gpt-4-turbo-preview</code>	OpenAI model
<code>LLM_TEMPERATURE</code>	0.0	Deterministic output
<code>LLM_MAX_TOKENS</code>	1024	Maximum response length
<code>LLM_TIMEOUT</code>	60s	Request timeout

## Confidence Scoring

The service calculates a confidence score based on:

Condition	Score
No answer found	0.0
1 relevant context	0.5-0.7
2-3 relevant contexts	0.7-0.85
4+ relevant contexts	0.85-0.95

## 3.3 Vector Store Service ( [core/vector\\_store.py](#) )

The Vector Store manages document embeddings using FAISS (Facebook AI Similarity Search).

### FAISS Index Configuration

```
Index Type: IndexFlatIP (Inner Product)
Dimension: 1536
Similarity: Cosine (vectors are normalized)
```

### Storage Architecture

```
./data/faiss_index/
├── faiss.index      # Binary FAISS index file
├── chunks.json      # Chunk text and metadata
└── documents.json   # Document-level metadata

./data/documents/
└── embedding_cache/ # Cached embeddings (NPY files)
```

## Key Operations

Operation	Description
<code>add_chunks()</code>	Add document chunks with embeddings
<code>similarity_search()</code>	Find top-k similar chunks
<code>delete_document()</code>	Remove document (rebuilds index)
<code>get_stats()</code>	Get index statistics


## Similarity Search Flow



## 3.4 Retriever Service ( `core/retriever.py` )

The Retriever orchestrates the complete RAG pipeline, connecting all components.

## Pipeline Flow



## Critical Design: Retrieve First

The retriever **ALWAYS** performs retrieval before generation. This ensures:

1. The LLM only sees relevant context
2. If no context is found, no LLM call is made
3. Resources are not wasted on unanswerable questions

# 4. Ingestion Pipeline

---

## Overview

---

The ingestion pipeline processes documents through three stages:



## Stage 1: Document Loading ([ingestion/loader.py](#))

---


### Supported Formats

Format	Loader	Features
<b>PDF</b>	<code>PDFLoader</code>	Page-by-page extraction with page numbers
<b>DOCX</b>	<code>DOCXLoader</code>	Section detection via heading styles
<b>Markdown</b>	<code>MarkdownLoader</code>	Header-based sections, HTML conversion

### Document Object

```
@dataclass
class Document:
    text: str
    metadata: dict[str, Any]
```

## Loader Selection



## Stage 2: Semantic Chunking ([ingestion/chunker.py](#))

### Chunking Strategy

The `SemanticChunker` splits documents intelligently:

1. **Split into sentences** using regex patterns
2. **Group sentences** to reach target chunk size
3. **Add overlap** from previous chunk for context continuity

### Configuration


Parameter	Default	Description
CHUNK_SIZE	512 tokens	Target chunk size
CHUNK_OVERLAP	50 tokens	Overlap between chunks

## Chunking Visualization

Original Document:

```
| Sentence 1. Sentence 2. Sentence 3. Sentence 4. Sentence 5. |
| Sentence 6. Sentence 7. Sentence 8. Sentence 9. Sentence 10. |
```

After Chunking (with overlap):




## Chunk Object

```
@dataclass
class Chunk:
    chunk_id: str          # Deterministic ID (SHA-256)
    text: str               # Chunk content
    metadata: dict[str, Any] # Source, page, section, index
```

## Stage 3: Ingestion Service ([ingestion/pipeline.py](#))

The `IngestionService` coordinates the complete flow:



# 5. Query Pipeline

## Complete Query Flow






## 6. Hallucination Prevention

This system implements **four layers** of hallucination prevention:

## Layer 1: Retrieval Gate


---



## Layer 2: System Prompt Constraint


---

The LLM receives strict instructions:




## Layer 3: Context Isolation

---



## Layer 4: Mandatory Citations

---



# 7. API Reference

---

## Base URL

---

```
http://localhost:8000
```

## Endpoints Overview

---

Method	Endpoint	Description
GET	/health	Health check with component status
POST	/ingest	Ingest documents
POST	/query	Ask questions

---

### GET /health

---

Check system health and component status.

#### Request

```
curl http://localhost:8000/health
```

## Response (200 OK)

```
{  
  "status": "healthy",  
  "version": "1.0.0",  
  "timestamp": "2024-12-25T10:30:00Z",  
  "components": {  
    "embeddings": "healthy",  
    "vector_store": "healthy",  
    "llm": "healthy"  
  }  
}
```

## POST /ingest

Ingest documents into the RAG system.

### Request Body

Field	Type	Required	Description
document_type	string	Yes	"pdf" , "docx" , or "markdown"
file_path	string	No*	Path to document file
content	string	No*	Raw markdown content
metadata	object	No	Custom metadata

\*Either `file_path` or `content` is required.

### Example: Ingest PDF

```
curl -X POST http://localhost:8000/ingest \  
-H "Content-Type: application/json" \  
-d '{  
  "document_type": "pdf",  
  "file_path": "C:/docs/policy.pdf",  
  "metadata": {"department": "HR"}  
}'
```

## Example: Ingest Markdown Content

```
curl -X POST http://localhost:8000/ingest \
-H "Content-Type: application/json" \
-d '{
  "document_type": "markdown",
  "content": "# Policy\n\nEmployees get 20 vacation days per year.",
  "metadata": {"source": "hr_policy"}
}'
```

## Response (200 OK)

```
{
  "success": true,
  "message": "Successfully ingested with 5 chunks",
  "documents": [
    {
      "document_id": "doc_a1b2c3d4e5f6",
      "source": "policy.pdf",
      "document_type": "pdf",
      "chunk_count": 5,
      "ingested_at": "2024-12-25T10:30:00Z",
      "metadata": {"department": "HR"}
    }
  ],
  "total_chunks": 5,
  "processing_time_ms": 1234.56
}
```

## POST /query

Ask questions about ingested documents.

### Request Body

Field	Type	Required	Default	Description
question	string	Yes	-	Question to ask
top_k	integer	No	5	Max documents to retrieve
similarity_threshold	float	No	0.7	Min relevance score
include_context	boolean	No	false	Include chunk text
metadata_filter	object	No	null	Filter by metadata

## Example Request

```
curl -X POST http://localhost:8000/query \
-H "Content-Type: application/json" \
-d '{
  "question": "How many vacation days do employees get?",
  "top_k": 3,
  "similarity_threshold": 0.8
}'
```

## Response (Answer Found)

```
{
  "success": true,
  "answer": "Employees get 20 vacation days per year. [Source 1]",
  "sources": [
    {
      "document_id": "doc_a1b2c3d4e5f6",
      "source": "hr_policy.pdf",
      "page_number": 5,
      "section": "Vacation Policy",
      "relevance_score": 0.89,
      "chunk_text": "Employees get 20 vacation days per year."
    }
  ],
  "confidence": 0.85,
  "query_time_ms": 1456.78,
  "retrieval_time_ms": 45.23,
  "generation_time_ms": 1411.55
}
```

## Response (No Answer)

```
{
  "success": true,
  "answer": "Answer not found in documents.",
  "sources": [],
  "confidence": 0.0,
  "query_time_ms": 45.23,
  "retrieval_time_ms": 45.23,
  "generation_time_ms": 0.0
}
```

---

# 8. Installation & Setup

---

## Prerequisites

---

Requirement	Version
Python	3.10+
pip	Latest
OpenAI API Key	Required

## Step-by-Step Installation

---

### Step 1: Clone/Navigate to Project

```
cd RAG
```

### Step 2: Create Virtual Environment

```
# Create virtual environment
python -m venv venv

# Activate (Windows)
venv\Scripts\activate

# Activate (Linux/Mac)
source venv/bin/activate
```

### Step 3: Install Dependencies

```
pip install -r requirements.txt
```

## Step 4: Configure Environment

```
# Copy example config  
cp .env.example .env  
  
# Edit .env file  
notepad .env # Windows  
# OR  
nano .env # Linux/Mac
```

Add your OpenAI API key:

```
OPENAI_API_KEY=sk-your-api-key-here
```

## Step 5: Run the Application

```
uvicorn app.main:app --reload --port 8000
```

## Step 6: Verify Installation

```
curl http://localhost:8000/health
```

Expected output:

```
{  
    "status": "healthy",  
    "components": {  
        "embeddings": "healthy",  
        "vector_store": "healthy",  
        "llm": "healthy"  
    }  
}
```

---

# 9. Configuration Reference

---

## All Environment Variables

---

Variable	Required	Default	Description
<strong>Application</strong>			
APP_NAME	No	Enterprise RAG System	Application name
APP_ENV	No	development	Environment (development/staging/production)
DEBUG	No	false	Enable debug mode
LOG_LEVEL	No	INFO	Logging level (DEBUG/INFO/WARNING/ERROR)
LOG_FORMAT	No	json	Log format (json/console)
<strong>API</strong>			
API_HOST	No	0.0.0.0	API server host
API_PORT	No	8000	API server port
CORS_ORIGINS	No	*	Allowed CORS origins
<strong>OpenAI</strong>			
OPENAI_API_KEY	Yes	-	OpenAI API key
OPENAI_ORG_ID	No	-	OpenAI organization ID
<strong>Embedding</strong>			
EMBEDDING_MODEL	No	text-embedding-3-small	OpenAI embedding model
EMBEDDING_DIMENSION	No	1536	Embedding vector dimension
EMBEDDING_BATCH_SIZE	No	100	Batch size for embeddings

<b>LLM</b>			
LLM_MODEL	No	gpt-4-turbo-preview	OpenAI LLM model
LLM_TEMPERATURE	No	0.0	Temperature (0=deterministic)
LLM_MAX_TOKENS	No	1024	Maximum response tokens
LLM_TIMEOUT	No	60	Request timeout (seconds)
<b>Retrieval</b>			
SIMILARITY_THRESHOLD	No	0.7	Minimum similarity score
TOP_K	No	5	Maximum documents to retrieve
<b>Chunking</b>			
CHUNK_SIZE	No	512	Chunk size (tokens)
CHUNK_OVERLAP	No	50	Overlap between chunks
<b>Storage</b>			
FAISS_INDEX_PATH	No	./data/faiss_index	FAISS index directory
DOCUMENT_STORE_PATH	No	./data/documents	Document metadata directory

---

# 10. Appendix

## A. Project Structure

```
RAG/
├── app/
│   ├── __init__.py
│   ├── config.py          # Pydantic Settings configuration
│   └── main.py            # FastAPI application entry point
|
├── api/
│   ├── __init__.py
│   ├── dependencies.py   # Dependency injection (singletons)
│   └── routes/
│       ├── __init__.py
│       ├── health.py      # GET /health endpoint
│       ├── ingest.py      # POST /ingest endpoint
│       └── query.py       # POST /query endpoint
|
├── core/
│   ├── __init__.py
│   ├── embeddings.py     # OpenAI embeddings + cache
│   ├── llm.py            # LLM service with anti-hallucination
│   ├── retriever.py      # RAG orchestration
│   └── vector_store.py   # FAISS vector store
|
├── ingestion/
│   ├── __init__.py
│   ├── loader.py          # Document loaders (PDF/DOCX/MD)
│   ├── chunker.py         # Semantic text chunking
│   └── pipeline.py        # Ingestion orchestration
|
└── schemas/
    ├── __init__.py
    ├── common.py           # HealthResponse, ErrorResponse
    ├── documents.py        # IngestRequest/Response
    └── query.py            # QueryRequest/Response
|
├── data/
│   ├── documents/         # Document metadata
│   │   └── embedding_cache/ # Cached embeddings
│   └── faiss_index/        # FAISS index persistence
│       ├── faiss.index
│       ├── chunks.json
│       └── documents.json
|
└── examples/
    ├── sample_policy.md    # Sample document
    └── test_rag.py          # Test script
|
├── .env.example           # Example environment config
└── requirements.txt        # Python dependencies
└── README.md              # Project readme
```

## B. Data Persistence

---

All data is persisted to the `./data/` directory:

File	Purpose
<code>faiss_index/faiss.index</code>	Binary FAISS index
<code>faiss_index/chunks.json</code>	Chunk text and metadata
<code>faiss_index/documents.json</code>	Document-level metadata
<code>documents/embedding_cache/</code>	Cached embedding vectors

**To reset all data:** Delete the `./data/` directory.

## C. Dependencies

---

### Core Framework

- `fastapi==0.109.0` - Web framework
- `uvicorn[standard]==0.27.0` - ASGI server
- `pydantic==2.5.3` - Data validation
- `pydantic-settings==2.1.0` - Settings management

### AI/ML

- `openai==1.12.0` - OpenAI API client
- `faiss-cpu==1.7.4` - Vector similarity search
- `numpy>=1.24.0,<2.0.0` - Numerical computing

### Document Processing

- `pypdf==4.0.1` - PDF text extraction
- `python-docx==1.1.0` - DOCX processing
- `markdown==3.5.2` - Markdown parsing
- `beautifulsoup4==4.12.3` - HTML parsing

## Utilities

- `python-dotenv==1.0.0` - Environment loading
- `tenacity==8.2.3` - Retry logic
- `httpx==0.26.0` - HTTP client
- `structlog==24.1.0` - Structured logging

## D. Troubleshooting

---

Issue	Solution
<code>OPENAI_API_KEY not set</code>	Add key to <code>.env</code> file
<code>Connection refused</code>	Check if server is running on port 8000
<code>No answer found</code>	Lower <code>similarity_threshold</code> or ingest more documents
<code>Rate limit exceeded</code>	Built-in retry handles this; wait and retry
<code>FAISS index corrupted</code>	Delete <code>./data/faiss_index/</code> and re-ingest

---

**End of Documentation**

---

*This documentation was generated for the Enterprise RAG System v1.0.0*